

TRACE-RPG: A Trace-Linked Symbolic Commit Gate for Generated Events in an Interactive Game World

Anonymous Author(s)

Abstract—Large language models can generate expressive quests and non-player-character dialogue, but fluent or schema-valid text does not establish that a proposed event is executable, authorized, or consistent with canonical game state. We present TRACE-RPG, a symbolic commit gate that isolates untrusted proposals from state mutation. A type-strict parser enforces the shared candidate-key contract, rejecting ambiguous known-field types and unknown top-level candidate fields; an externally supplied action policy and deterministic state-relative predicates then decide every candidate transition admitted by this controller. Rejected candidates are exposed to a bounded repair callback, and successful candidates are defensively revalidated before mutation. Terminal experiment records are linked to unkeyed SHA-256 content checksums, state-semantic replay recomputes the recorded transition, and episode-continuity checks reject disconnected histories. We evaluate the implementation on purposefully designed deterministic offline fixtures over a single authored world state, and we mark which reported quantities are observations and which are construction invariants of the harness. The reference repair callback reconstructs policy fields from the authoritative state, so it is an oracle upper bound rather than a deployable repair method, and the adapter telemetry values are synthetic fixture constants rather than measurements. The evidence establishes implementation conformance for declared structured fields on one world, and it establishes nothing about live-model behaviour, player experience, affect, retrieval, memory, or commercial-engine performance.

Index Terms—game artificial intelligence, interactive narrative, symbolic validation, constrained generation, runtime enforcement, reproducibility

I. INTRODUCTION

INTERACTIVE generation is more than a language task. A candidate event may be natural yet require an absent item, assert a fact unknown to its speaker, reveal forbidden quest information, introduce an unauthorized effect, or regress a quest stage. Executable text environments make the distinction between a plausible utterance and a valid transition explicit [1], [2]. Grounded dialogue and personalized quest systems further show the value of quest, entity, and graph context [3], [4], but context does not itself authorize canonical mutation.

Structural generation constraints solve an adjacent problem. Grammar-constrained decoding, language-model query languages, and incremental parsers can restrict generated structures [5], [6], [7]. A syntactically admissible event can nevertheless be false relative to the current world. Retrieval and memory can supply useful evidence [8], [9], [10], but retrieved content is also not a state-transition authority.

This paper presents TRACE-RPG, an implementation that combines these concerns as an auditable transaction protocol. Its current contributions are:

- 1) versioned contracts for canonical state, externally supplied action policy, candidate event, experiment record, and game bridge;
- 2) deterministic state-relative validation, a bounded repair callback with unchanged fallback, and defensive pre-commit validation;
- 3) content-linked terminal records, state-semantic replay, and episode-continuity validation; and
- 4) an offline harness with type-strict known-field adapter boundaries, frozen assignment keys, distinct terminal failure classes, and descriptive treatment-policy accounting.

The empirical center is a deterministic designed-fixture conformance pilot over a single authored world state; a retained authored Godot 4.7.1 policy-mirror trace and render replay add bounded engine-local evidence. Headless runs of that fixture do record engine-local timing, and we report it as such: each run is one scripted fixture at a fixed seed, and its per-request figure measures local fixture handling rather than model or network inference. Across the four fixture runs, the largest sampled headless frame delta was the first of five samples in every run (98.760–116.667 ms), so the 16.7 ms frame budget was not met; these startup-sensitive traces are not a stable frame benchmark. There are no live model calls, participants, affect or retrieval experiments, live Python–Godot authorization round trips, or representative end-to-end gameplay performance measurements. Those studies are outside this paper.

II. RELATED WORK AND RESEARCH GAP

A. Interactive Narrative and Game Agents

KNUDGE evaluates branching NPC dialogue grounded in quest and entity specifications [3]. Knowledge-graph-assisted quest generation [4], memory–reflection–planning agents [11], and trainable role-playing agents [12] address complementary aspects of grounded and believable behavior. Player-facing studies evaluate preference and perceived dialogue quality [13], [14]. Yin *et al.*’s 2026 online record has an issue assignment that was future-dated at our metadata audit and must be rechecked at submission [15]. A 2025 preprint reports role-sensitive stability–improvisation trade-offs [16]; we use it only as recent motivation. These works motivate separating hard

admissibility from optional naturalness and player-experience constructs.

Experience-driven content generation [17] and observer-labelled engagement [18] motivate a future adaptation track. A recent preprint reports limitations in vision-language-model engagement inference [19]. Accordingly, affect remains non-authoritative and is absent from the present implementation evidence.

B. Environments and Benchmarks

TextWorld and ScienceWorld expose explicit interactive state [1], [2]; MineDojo adds open-ended, knowledge-rich embodied tasks [20]. General Video Game AI separates agent, game, and content tracks [21]. AgentBench studies agents across heterogeneous environments [22], AgentBoard adds process-level diagnostics [23], SOTOPIA evaluates social interaction [24], and BALROG evaluates long-horizon language and visual game reasoning [25]. Automated play-style agents provide a complementary testing precedent, not a substitute for human evaluation [26]. These benchmarks inform our future evaluation design but do not prescribe a transactional authorization boundary.

C. Constraints, Neuro-Symbolic Validation, and Retrieval

Constrained decoders improve output structure [5], [6], [7]; TRACE-RPG instead treats structure as an input condition to independent state-relative authorization. Two direct game-specific comparators span complementary boundaries. PANGeA combines memory, validation, a REST service, and Unity integration, but its documented validation asks an LLM to judge and iteratively refine content against designer rules [27]. IVIE, currently a preprint associated with non-archival ICCG 2026 material, incrementally generates interactive-fiction worlds with symbolic validation [28]. Neither comparison implies superiority: each reinforces that validity is only validity for the objectives and predicates its system encodes.

G-Retriever, HippoRAG, and Mem0 address relevant-subgraph retrieval and long-term memory [8], [9], [10]; Generative Agents and Character-LLM address behavioral memory and persona [11], [12]. TRACE-RPG may later consume such context, but denies it direct commit authority. Retrieval and temporal-memory benefits are not measured here.

D. Auditability and Reproducibility

Process metrics motivate retaining intermediate failures [23]; robust evaluation cautions against overinterpreting sparse stochastic runs [29]. Reproducibility and environmental reporting motivate frozen artifacts and explicit measurement boundaries [30], [31]. Human and LLM evaluation require construct, rater, and bias controls [32], [33], [34], [35], while future non-inferiority and participant analyses require justified margins and random-effects structures [36], [37].

The bounded novelty is the implemented combination, not the invention of its individual parts. Action interposition is long established. Narrative mediation intercepts a user or agent action before it executes and either blocks it or accommodates it by revising the plan [38]; classical planning

defines applicability through preconditions and effects [39]; runtime enforcement wraps an untrusted actor in a monitor that admits only permitted transitions [40]; authored-logic interactive drama governs which utterances a character may make, in this journal’s predecessor title [41]; and belief-aware narrative planning reasons over what characters know, which is the concern our knowledge and disclosure predicates encode far more coarsely [42]. TRACE-RPG does not claim to have invented interposition. Its narrower contribution is the evidence layer bound to that gate under an untrusted generative proposer: terminal records linked to content checksums, state-semantic replay of the recorded transition, episode-continuity validation, and assignment-complete failure accounting that retains cases for which no proposal trace can exist. We do not infer that a feature absent from a report is absent from every implementation.

III. PROBLEM DEFINITION AND GUARANTEE BOUNDARY

For exposition, the controller input used by validation at step t is decomposed as

$$c_t = (G_t, q_t), \quad (1)$$

where G_t denotes typed snapshot fields and q_t denotes the externally authored action-policy, disclosure, and quest-stage constraints stored with that snapshot. Terminal record history $h_{<t}$ is maintained separately from the `WorldState` snapshot; no validation predicate reads it, and the implementation does not reconstruct state from the log. Retrieval, narrative scores, model rationales, memory summaries, and any future affect estimate are non-authoritative proposal context.

An untrusted proposer emits candidate a_t . Strict parsing first enforces the shared candidate-key contract, including rejection of unknown top-level fields, and establishes the declared schema. The validator returns

$$V(c_t, a_t) = (v_{\text{policy}}, v_{\text{pre}}, v_{\text{reach}}, v_{\text{know}}, v_{\text{disc}}, v_{\text{quest}}, E_t), \quad (2)$$

where E_t is a structured error set. Mutation is

$$c_{t+1} = \begin{cases} T(c_t, a_t), & \bigwedge_i v_i = 1, \\ c_t, & \text{otherwise.} \end{cases} \quad (3)$$

The controller asserts the following implementation invariants under its encoded contracts; the pilot tests their mechanisms with deterministic valid, invalid, repair, fallback, replay, and fault-injection fixtures rather than presenting general theorems. **I1 (exhausted-failure immutability):** if no returned candidate through budget K validates, or proposal/controller execution fails before commit, the returned state equals the supplied prior state. **I2 (bounded recorded attempts):** provided every proposal and repair callback returns, repair budget $K \geq 0$ admits at most $K + 1$ recorded candidate attempts; a successful commit adds one defensive validation immediately before mutation. The runtime does not enforce callback deadlines, so I2 is not a wall-clock termination guarantee. **I3 (effect authorization):** an accepted candidate contains no

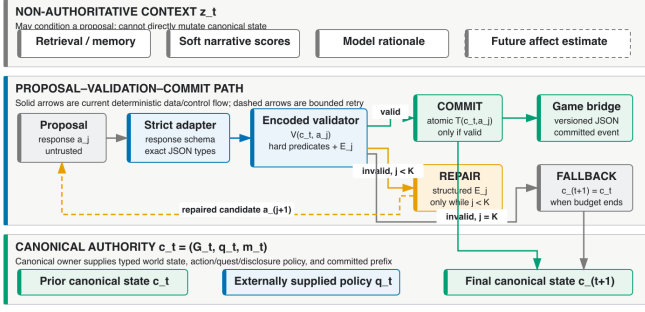


Fig. 1. Trust boundary. Retrieved context and generated candidates remain non-authoritative until known-field type checks and state-relative validation complete.

TABLE I
ENCODED VALIDATOR BOUNDARY

Predicate	Rejects	On failure
Action policy	unknown action/type	unchanged
Precondition	absent/false requirement	unchanged
Reachability	inaccessible object	unchanged
NPC knowledge	undeclared known fact	unchanged
Disclosure	forbidden fact	unchanged
Quest	ineligible/regressive stage	unchanged

declared fact effect in `CandidateAction.effects` outside `ActionPolicy.allowed_effects`, and any quest-stage target is explicitly authorized by that policy. **I4 (state-semantic replay)**: with the frozen schemas and deterministic transitions, an accepted trace reconstructs its stated terminal state.

These asserted invariants assume correct policy authorship, complete and faithful extraction into declared candidate fields, deterministic software versions, returning callbacks, and single-process file ownership. State-semantic replay revalidates the recorded structured candidate and transition; it does *not* reproduce or verify the provenance, reasoning, prompts, or stochastic generation of intermediate repairs. Natural-language implications omitted from structured fields remain outside the guarantee.

IV. TRACE-RPG ARCHITECTURE AND ALGORITHMS

A. Trust Boundary and Contracts

Fig. 1 places the model or recorded adapter outside the canonical owner. The parser rejects ambiguous numeric and collection types for known fields and rejects unknown top-level candidate fields under the shared key contract. The action policy is supplied separately from the candidate and defines required preconditions and allowed effects. The game bridge is a versioned event interface that may carry observations, proposals, validations, rejections, fallbacks, or commits; only a valid commit authorizes canonical mutation. This decouples interfaces but does not yet demonstrate portability to a commercial engine. Table I summarizes the encoded rejection boundary; it is not a claim of semantic completeness.

```

1:  $a \leftarrow \text{PROPOSE}(c); j \leftarrow 0$ 
2: repeat
3:    $(ok, E) \leftarrow V(c, a); \text{RECORD}(a, E)$ 
4:   if  $ok$  then
5:     assert  $V(c, a).ok$ ; return  $\text{COMMIT}(c, a)$ 
6:   if  $j = K$  then return  $\text{FALLBACK}(c)$ 
7:    $a \leftarrow \text{REPAIR}(c, a, E); j \leftarrow j + 1$ 
8: until terminal

```

Fig. 2. Validate-repair-commit procedure. At most $K + 1$ candidate attempts are recorded; the successful path performs a defensive validation before mutation.




Fig. 3. Proposal, validation, bounded repair, fallback, commit, and replay states. Controller failure is terminal and may have no completed proposal trace.

B. Validate-Repair-Commit

Fig. 2 gives the controller’s dependency-free pseudocode. “Record” denotes an append to the in-memory trace representation; persistent writing occurs only after terminal consistency checks.

Validation errors are exposed to a repair callback as structured error codes. Repair never directly mutates state, and exhaustion returns an unchanged deterministic fallback. Fig. 3 shows terminal paths. The reference callback used in the pilot is deliberately *not* counterexample-guided: it discards the error set and reconstructs the policy-required preconditions and effects directly from the authoritative state. It is therefore an oracle upper bound that exercises control flow, and its outcomes must not be read as evidence about any repair method that consumes counterexamples.

C. Integrity, Semantic Replay, and Accounting

Each result row carries an unkeyed SHA-256 *integrity checksum*; proposal outcomes additionally link to a trace checksum. These hashes detect accidental or tested content changes but are not signatures, message-authentication codes, writer identity, or protection against an actor able to rewrite content and checksums. Semantic replay goes beyond byte integrity by strictly parsing the prior state and every recorded candidate/validation pair, requiring every nonfinal attempt to be invalid, recomputing the terminal transition, and comparing the stated final state. Episode replay also requires state continuity between adjacent JSONL records. Replay does not authenticate or re-execute the repair-generation step that produced candidate $j + 1$.

The terminal record classes are commit, symbolic fallback, adapter failure, and controller failure. A controller-failure row remains in the frozen assignment denominator and leaves state unchanged, but it may legitimately lack a completed proposal trace because no candidate outcome was produced. Accordingly, result completeness does not imply trace existence for every terminal row. The pre-execution assignment key contains run, arm, scenario, seed, model, revision, controller-configuration hash, assignment-input hash, and prior-state hash. Aggregation rejects duplicate, missing, or unexpected keys and separately reports symbolic, adapter, and controller failures.

V. DETERMINISTIC OFFLINE PILOT METHOD

A. Questions and Designed Fixtures

The pilot asks whether: (PQ1) valid candidates commit while named invalid-policy cases remain unchanged; (PQ2) rejection-only, unchanged retry, and structured repair produce the expected outcomes on authored repairable and irreparable cases; (PQ3) frozen checksum, state-semantic, linkage, and continuity mutations are rejected by their designated check operations, while an injected second-file append failure triggers the expected pair-write rollback check; (PQ4) strict adapters and assignment manifests retain exactly one classified terminal row per assigned case or fail closed; and (PQ5) proposal and replay parsers both reject a complete otherwise-valid candidate carrying one unknown top-level key.

Cases are purposively designed conformance fixtures, not samples from a model or game population. The four domains are validator/state isolation, repair control flow, record/trace fault injection, and adapter/accounting contracts. Unique cases, rather than repeated deterministic seeds, form descriptive denominators. All expectations are specified before execution.

B. Measures and Reproducibility

Measures are exact counts: expected/observed validator-class agreement; rejected or controller-failure cases with state mutation; commit and attempt outcomes by repair arm and reparability; rejected/prespecified detectable faults by designated operation; replayed/committed traces; rejected/injected manifest violations; and proposal/replay rejection agreement on the prespecified unknown-key negative regression. The last measure is deterministic parser-contract parity, not an empirical safety outcome. Exception class alone does not identify a stable detector layer, so no detector-attribution result is reported. No p -values or population confidence intervals are attached to these designed fixtures. The adapter cases carry synthetic latency and token constants that the input schema pins with JSON-Schema `const`; they exercise accounting-field propagation only and are not measurements of any provider, runtime, or device.

The release packet freezes fixture and assignment manifests, schemas, a result JSON artifact with embedded outcome and trace records, environment and code revision, generated tables, and checksums. Reporting follows reproducibility and measurement-boundary guidance [30], [31].

VI. PILOT RESULTS

A. Declared-Field Gate Conformance

The pilot loader admits a fixture set only when it isolates every implemented validator code exactly once alongside one valid control, so the 13/13 agreement and the observation of all 12 implemented error codes are *construction invariants of the harness*, not measured success rates. What the run adds is the part the loader does not enforce: each of the 12 isolated negative fixtures showed observed agreement with its authored expected code rather than reaching a different one, and every noncommit path preserved the complete canonical state. This is implementation conformance to an authored structured-field oracle over a single world state, not accuracy against independent semantic labels.

B. Repair Mechanism

Across two initially invalid cases, rejection-only and unchanged-candidate retry each committed 0/2. The reference repair callback committed the repairable precondition case after one repair and left the deliberately nonrepairable reachability case at fallback, for 1/2 commits. That callback is not counterexample-guided: it discards the error set and copies the policy-required preconditions and effects from the authoritative state, so it is an oracle upper bound whose outcome is determined by which candidate fields it assigns. The three arms are therefore a control-flow trace, not a comparison of repair strategies.

C. Integrity, Boundaries, and Assignment Accounting

All 10 faults prespecified as detectable were rejected by their designated check operations (10/10). Because this pilot does not compare stable typed detector codes, it does not establish detector-layer attribution. Separately, the known provenance-boundary fixture substituted a different invalid precursor before the same recorded valid repair, recomputed the checksum, and passed replay as expected (1/1 undetected). Thus, the mechanism does not protect against a writer who can recompute unkeyed checksums, and replay does not authenticate the repair-generation operation. Two open boundary sentinels—unextracted narrative disclosure and simultaneous omission of a required object from candidate and policy—were intentionally accepted at the encoded layer (2/2), with 0 labelled as safety passes. The former unknown-top-level-field acceptance boundary was reclassified after Stage 8 as a closed negative regression: both proposal and replay parsers specifically rejected a complete 12-field candidate carrying one unknown key (1/1). This establishes parser rejection parity, not semantic safety. Among seven adapter/accounting assignments, outcomes were 1 commit, 1 symbolic fallback, and 5 adapter failures. Synthetic telemetry fields were populated and propagated for 2/7 assignments; these are schema-pinned constants verifying accounting-field propagation, not measurements. All 3/3 injected duplicate or missing-assignment guards failed closed. These are raw counts from synthetic offline fixtures.



Panel	Capture	Event	Validation	State relation
(a)	sl-rc-001-arrival	evt-000-observe	not applicable	snapshot = before
(b)	sl-rc-002-rejected-secret	evt-002-fallback-secret	invalid	world = before
(c)	sl-rc-003-authorized-hint	evt-005-commit-hint	valid	world $\neq$ before

Full-resolution artifact pointer: godot-4.7.1-20260813t115916z-sealed-lighthouse-render-v5/rendered-canonical-v1/; filenames are the capture IDs above plus .png.

Fig. 4. Illustrative status/event crops from the authored Godot render replay, with manifest metadata. The trace-bound correspondence is not live-transport, performance, usability, immersion, or efficacy evidence.

TABLE II
REPAIR-CALLBACK CONTROL FLOW (NOT A COMPARISON)

Callback	Injected defect	Fields it assigns	Terminal path
None (K=0)	unreachable object (nonrepairable)	—	fallback
None (K=0)	policy precondition (repairable)	—	fallback
Unchanged retry	unreachable object (nonrepairable)	none	fallback
Unchanged retry	policy precondition (repairable)	none	fallback
Policy-restore oracle	unreachable object (nonrepairable)	preconditions, effects	fallback
Policy-restore oracle	policy precondition (repairable)	preconditions, effects	commit

TABLE III
FROZEN-PILOT RAW ACCOUNTING

Check	Numerator	Denominator
Gate fixture agreement ^c	13	13
Known-boundary acceptances ^a	2	2
Closed unknown-key boundary rejected ^d	1	1
Detectable integrity faults rejected	10	10
Known integrity boundary replay-accepted ^b	1	1
Adapter commits	1	7
Symbolic fallbacks	1	7
Adapter failures	5	7
Manifest guards rejected	3	3

^aBoundary acceptances document semantic-extraction and policy-completeness limits; they are not safety passes. ^bReplay does not authenticate or re-execute repair generation. ^cThe loader enforces this agreement, so it is a construction invariant rather than a measurement. ^dThis is post-Stage-8 parser rejection parity, not semantic-safety evidence.

Tables II and III separate the authored callback paths from raw fixture accounting so neither is read as a strategy comparison or population estimate.

Fig. 5 makes the reporting boundary explicit: only rows admitted by the frozen fixture and designated check operation enter the generated result tables.

Godot 4.7.1 replayed the retained *Sealed Lighthouse* canonical trace in a non-headless 1280×720 SubViewport. Its manifest binds each PNG to event/delivery indices, validation, and state hashes: Fig. 4(b) preserves the before hash, whereas (c) changes it after a valid commit. All source panels are generated-asset-free primary-track renders; the illustrative




Fig. 5. Evidence boundary from frozen fixture through a designated check operation and generated table. Counts are admitted only through generated artifacts.

crops below enlarge only their authored status/event callouts.

The selected v5 packet establishes one scripted fixture’s render/state correspondence only; it adds no live model or Python transport, participant or player input, adaptation intervention, or narrative-quality result.

VII. DISCUSSION

A. Established Scope

The designed pilot can establish mechanism behavior only for its frozen cases: whether encoded checks isolate state, whether bounded repair and fallback follow their declared paths, whether named corruptions are rejected by their designated operations, and whether assigned rows remain classified for aggregation. It does not attribute a rejection to a stable typed detector code. This is useful for authoring-time validation, continuous-integration regression, recorded-adapter tests, and pre-engine replay inspection.

The mechanism does not establish that the policy is correctly authored, that natural-language facts have been completely extracted, or that every semantically unsafe event is represented by a declared field. Nor do unkeyed hashes authenticate the writer. The writer assumes single-process ownership; concurrent-writer safety is not demonstrated. A versioned bridge is interface decoupling, not evidence of engine portability or real-time performance.

B. Relationship to Prior Work

TRACE-RPG complements constrained decoders by applying state-relative authorization after parsing [5], [6], [7]. It

differs from retrieval and memory components by withholding commit authority [8], [9], [10]. It complements executable environments and benchmarks rather than replacing them [1], [2], [22], [25]. Process-linked traces instantiate a measurement concern emphasized by AgentBoard [23]; they are not evidence that TRACE-RPG outperforms that or any other system.

The studies that would test efficacy—a preregistered multi-model comparison with an independently labelled semantic oracle and a genuine blind-retry comparator, retrieval and memory ablations at matched encoded validity, and a blinded matched-validity human study with justified power and mixed-effects analysis—are outside this paper and are not reported here in any partial form [22], [32], [29], [37].

VIII. THREATS, ETHICS, AND ARTIFACT SCOPE

Construct validity is bounded by declared fields; an encoded check cannot discover omitted semantics. Internal validity depends on independently authored policy expectations and absence of fixture leakage. External validity is limited by offline text/mock cases. Conclusion validity is protected by reporting raw designed-case counts without population inference. Reproducibility depends on fixed software and artifact versions. Security validity is restricted to unkeyed content-integrity checks and deterministic semantic replay, not adversarial authentication.

The pilot processes no participants or personal telemetry. Future human and affect work requires applicable review, informed consent, compensation disclosure, minimization, retention/deletion rules, and opt-out. LLM judges, if used, remain auxiliary because both automated and crowd judgements require calibration and bias controls [32], [33], [34], [35].

IX. CONCLUSION

TRACE-RPG isolates authored candidate-event payloads from canonical mutation through an encoded symbolic commit gate and makes terminal outcomes auditable through linked records and state-semantic replay. The deterministic offline pilot is evidence about the implemented mechanism in frozen fixture domains over a single authored world state, not cross-model superiority, player-experience benefit, retrieval or memory effectiveness, affect efficacy, or commercial-engine performance. These remain explicit confirmatory extensions.

DISCLOSURE OF AI USAGE

Large language models assisted prose, code, and test drafting; command orchestration; evidence auditing and result interpretation; the reviewer-panel simulation that drove this revision; and citation-identity checks against bibliographic metadata APIs. They did not replace deterministic execution as the source of reported values. Pilot counts are emitted by the project runner from authored fixtures and enter the manuscript through machine-generated fragments; Godot timing and render statements are transcribed from the retained engine-evidence packet. Implementation and result claims in this revision were cross-checked against source and artifacts, and every citation identity was checked against at least one independent index. The authors take full responsibility for the content.

DATA AND CODE AVAILABILITY

The repository contains the implementation, pilot bundle, fixture and assignment manifests, generated result fragments, and build sources. The current Stage-10 SHA-256 manifest covers 35 artifacts and 21 declared inputs and binds 64 provenance rows (43 executed fixture rows and 21 aggregate rows). It records `dirty=false` at input commit `c4752df`, tagged `trace-rpg-stage10-inputs-20260814-v1`. All 35 artifact hashes and all 21 input hashes recompute, and every input path and exact digest matches that commit (56/56 total hashes). The manifest records the portable invocation `uv run python` and contains no absolute user or clone path. This closes the clean-input reproducibility gate; it does not mean that an anonymized reviewer archive or DOI-bearing deposit already exists, and neither is claimed here.

REFERENCES

- [1] M.-A. Côté, Á. Kádár, X. Yuan *et al.*, “TextWorld: A learning environment for text-based games,” in *Computer Games*, ser. Communications in Computer and Information Science. Springer International Publishing, 2019, vol. 1017, pp. 41–75.
- [2] R. Wang, P. Jansen, M.-A. Côté, and P. Ammanabrolu, “ScienceWorld: Is your agent smarter than a 5th grader?” in *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2022, pp. 11 279–11 298.
- [3] N. Weir, R. Thomas, R. d’Amore, K. Hill, B. Van Durme, and H. Jhamtani, “Ontologically faithful generation of non-player character dialogues,” in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2024, pp. 9212–9242.
- [4] T. Ashby, B. K. Webb, G. Knapp, J. Searle, and N. Fulda, “Personalized quest and dialogue generation in role-playing games: A knowledge graph- and language model-based approach,” in *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*. ACM, 2023, pp. 1–20.
- [5] S. Geng, M. Josifoski, M. Peyrard, and R. West, “Grammar-constrained decoding for structured NLP tasks without finetuning,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2023, pp. 10932–10952.
- [6] L. Beurer-Kellner, M. Fischer, and M. Vechev, “Prompting is programming: A query language for large language models,” *Proceedings of the ACM on Programming Languages*, vol. 7, no. PLDI, pp. 1946–1969, 2023.
- [7] T. Scholak, N. Schucher, and D. Bahdanau, “PICARD: Parsing incrementally for constrained auto-regressive decoding from language models,” in *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2021, pp. 9895–9901.
- [8] X. He, Y. Tian, Y. Sun *et al.*, “G-Retriever: Retrieval-augmented generation for textual graph understanding and question answering,” in *Advances in Neural Information Processing Systems*, vol. 37. Neural Information Processing Systems Foundation, 2024, pp. 132 876–132 907.
- [9] B. Jiménez Gutiérrez, Y. Shu, Y. Gu, M. Yasunaga, and Y. Su, “HippoRAG: Neurobiologically inspired long-term memory for large language models,” in *Advances in Neural Information Processing Systems*, vol. 37. Neural Information Processing Systems Foundation, 2024, pp. 59 532–59 569.
- [10] P. Chhikara, D. Khant, S. Aryan, T. Singh, and D. Yadav, “Mem0: Building production-ready AI agents with scalable long-term memory,” in *ECAI 2025*, ser. Frontiers in Artificial Intelligence and Applications. IOS Press, 2025, peer-reviewed ECAI 2025 record; page range was not exposed by the accessible Crossref or publisher metadata at audit time.
- [11] J. S. Park, J. C. O’Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein, “Generative agents: Interactive simulacra of human behavior,” in *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*. ACM, 2023, pp. 1–22.
- [12] Y. Shao, L. Li, J. Dai, and X. Qiu, “Character-LLM: A trainable agent for role-playing,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2023, pp. 13 153–13 187.
- [13] N. Akoury, Q. Yang, and M. Iyyer, “A framework for exploring player perceptions of LLM-generated dialogue in commercial video games,” in *Findings of the Association for Computational Linguistics: EMNLP 2023*. Association for Computational Linguistics, 2023, pp. 2295–2311.
- [14] M. Hochreiter, S. Kriglstein, and G. Wallner, “Beyond pre-defined scripts: Player perceptions on generative non-player character dialogues,” in *Proceedings of the 31st International Conference on Intelligent User Interfaces*. ACM, 2026, pp. 2004–2018.
- [15] M. Yin, H. Zu, W. Gu *et al.*, “How contextualized generative AI shapes player experience in games,” *Entertainment Computing*, vol. 58, p. 101194, 2026, online record with a future issue assignment at the 2026-08-12 audit; issue metadata must be rechecked at submission.
- [16] V. Figueiredo and D. Elumeze, “Symbolically scaffolded play: Designing role-sensitive prompts for generative NPC dialogue,” 2025, preprint; no archival venue or publisher DOI verified as of 2026-08-12.

- [17] G. N. Yannakakis and J. Togelius, "Experience-driven procedural content generation," *IEEE Transactions on Affective Computing*, vol. 2, no. 3, pp. 147–161, 2011.
- [18] D. Melhart, D. Gravina, and G. N. Yannakakis, "Moment-to-moment engagement prediction through the eyes of the observer: PUBG streaming on Twitch," in *International Conference on the Foundations of Digital Games*. ACM, 2020, pp. 1–10.
- [19] Z. Wang, Q. Guo, R. Singh, and X. Hu, "Do vision language models understand human engagement in games?" 2026, preprint; no archival venue or publisher DOI verified as of 2026-08-12.
- [20] L. Fan, G. Wang, Y. Jiang *et al.*, "MineDojo: Building open-ended embodied agents with internet-scale knowledge," in *Advances in Neural Information Processing Systems*, vol. 35. Neural Information Processing Systems Foundation, 2022, pp. 18 343–18 362.
- [21] D. Pérez-Liébana, J. Liu, A. Khalifa, R. D. Gaina, J. Togelius, and S. M. Lucas, "General video game AI: A multitask framework for evaluating agents, games, and content generation algorithms," *IEEE Transactions on Games*, vol. 11, no. 3, pp. 195–214, 2019.
- [22] X. Liu, H. Yu, H. Zhang *et al.*, "AgentBench: Evaluating LLMs as agents," in *International Conference on Learning Representations*, 2024, pp. 52 989–53 046, peer-reviewed official proceedings record; no publisher DOI assigned.
- [23] C. Ma, J. Zhang, Z. Zhu *et al.*, "AgentBoard: An analytical evaluation board of multi-turn LLM agents," in *Advances in Neural Information Processing Systems*, vol. 37. Neural Information Processing Systems Foundation, 2024, pp. 74 325–74 362.
- [24] X. Zhou, H. Zhu, L. Mathur *et al.*, "SOTOPIA: Interactive evaluation for social intelligence in language agents," in *International Conference on Learning Representations*, 2024, pp. 40 975–41 019, peer-reviewed official proceedings record; no publisher DOI assigned.
- [25] D. Paglieri, B. Cupiał, S. Coward *et al.*, "BALROG: Benchmarking agentic LLM and VLM reasoning on games," in *International Conference on Learning Representations*, 2025, pp. 96 666–96 702, peer-reviewed official proceedings record; no publisher DOI assigned.
- [26] P. Le Pelletier de Woillemont, R. Labory, and V. Corruble, "Automated play-testing through RL-based human-like play-styles generation," *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, vol. 18, no. 1, pp. 146–154, 2022.
- [27] S. Buongiorno, L. Klinkert, Z. Zhuang, T. Chawla, and C. Clark, "PANGeA: Procedural artificial narrative using generative AI for turn-based, role-playing video games," *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, vol. 20, no. 1, pp. 156–166, 2024.
- [28] M. Vaucher, S. Silveira, S. Góngora, and L. Chiruzzo, "IVIE: A neuro-symbolic approach to incremental and validated generation of interactive fiction worlds," 2026, preprint; accepted/forthcoming at ICCV 2026, whose available pre-proceedings were explicitly non-archival and unpublished at retrieval time.
- [29] R. Agarwal, M. Schwarzer, P. S. Castro, A. Courville, and M. G. Bellemare, "Deep reinforcement learning at the edge of the statistical precipice," in *Advances in Neural Information Processing Systems*, vol. 34. Curran Associates, Inc., 2021, pp. 29 304–29 320, peer-reviewed official proceedings record; no publisher DOI assigned.
- [30] J. Pineau, P. Vincent-Lamarre, K. Sinha *et al.*, "Improving reproducibility in machine learning research: A report from the NeurIPS 2019 reproducibility program," *Journal of Machine Learning Research*, vol. 22, no. 164, pp. 1–20, 2021, peer-reviewed official journal article; no publisher DOI assigned.
- [31] P. Henderson, J. Hu, J. Romoff, E. Brunskill, D. Jurafsky, and J. Pineau, "Towards the systematic reporting of the energy and carbon footprints of machine learning," *Journal of Machine Learning Research*, vol. 21, no. 248, pp. 1–43, 2020, peer-reviewed official journal article; no publisher DOI assigned.
- [32] C. van der Lee, A. Gatt, E. van Miltenburg, S. Wubben, and E. Krahmer, "Best practices for the human evaluation of automatically generated text," in *Proceedings of the 12th International Conference on Natural Language Generation*. Association for Computational Linguistics, 2019, pp. 355–368.
- [33] M. Karpinska, N. Akoury, and M. Iyyer, "The perils of using Mechanical Turk to evaluate open-ended text generation," in *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2021, pp. 1265–1285.
- [34] L. Zheng, W.-L. Chiang, Y. Sheng *et al.*, "Judging LLM-as-a-judge with MT-Bench and Chatbot Arena," in *Advances in Neural Information Processing Systems*, vol. 36. Neural Information Processing Systems Foundation, 2023, pp. 46 595–46 623.
- [35] G. H. Chen, S. Chen, Z. Liu, F. Jiang, and B. Wang, "Humans or LLMs as the judge? a study on judgement bias," in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2024, pp. 8301–8327.
- [36] D. Lakens, A. M. Scheel, and P. M. Isager, "Equivalence testing for psychological research: A tutorial," *Advances in Methods and Practices in Psychological Science*, vol. 1, no. 2, pp. 259–269, 2018.
- [37] D. J. Barr, R. Levy, C. Scheepers, and H. J. Tily, "Random effects structure for confirmatory hypothesis testing: Keep it maximal," *Journal of Memory and Language*, vol. 68, no. 3, pp. 255–278, 2013.
- [38] M. O. Riedl, C. J. Saretto, and R. M. Young, "Managing interaction between users and agents in a multi-agent storytelling environment," in *Proceedings of the Second International Joint Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, 2003, pp. 741–748.
- [39] R. E. Fikes and N. J. Nilsson, "STRIPS: A new approach to the application of theorem proving to problem solving," *Artificial Intelligence*, vol. 2, no. 3–4, pp. 189–208, 1971.
- [40] F. B. Schneider, "Enforceable security policies," *ACM Transactions on Information and System Security*, vol. 3, no. 1, pp. 30–50, 2000.
- [41] R. Evans and E. Short, "Versu—a simulationist storytelling system," *IEEE Transactions on Computational Intelligence and AI in Games*, vol. 6, no. 2, pp. 113–130, 2014, predecessor title of IEEE Transactions on Games.
- [42] S. G. Ware and C. Siler, "Sabre: A narrative planner supporting intention and deep theory of mind," in *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment (AIIDE)*, vol. 17, no. 1, 2021, pp. 99–106.