

TRACE-RPG: A Trace-Linked Symbolic Commit Gate for Generated Events in an Interactive Game World

Anonymous Author(s)

Abstract—Large language models can generate expressive quests and non-player-character dialogue, but fluent or schema-valid text does not establish that a proposed event is executable, authorized, or consistent with canonical game state. We present TRACE-RPG, a symbolic commit gate that isolates untrusted proposals from state mutation. A type-strict parser enforces the shared candidate-key contract, rejecting ambiguous known-field types and unknown top-level candidate fields; an externally supplied action policy and deterministic state-relative predicates then decide every candidate transition admitted by this controller. Rejected candidates are exposed to a bounded repair callback, and successful candidates are defensively revalidated before mutation. Terminal experiment records are linked to unkeyed SHA-256 content checksums, state-semantic replay recomputes the recorded transition, and episode-continuity checks reject disconnected histories. We evaluate the implementation on purposefully designed deterministic offline fixtures over a single authored world state, and we mark which reported quantities are observations and which are construction invariants of the harness. A counterexample-guided repair operator ρ consumes only the prior candidate and the validator’s structured error set—never the authoritative state—and is compared against blind unchanged retry and a state-reading reference callback in a deterministic designed-fixture comparison over frozen repairability classes; the state-reading callback remains an oracle upper bound rather than a deployable repair method, and the adapter telemetry values are synthetic fixture constants rather than measurements. The evidence establishes implementation conformance for declared structured fields on one world, and it establishes nothing about live-model behaviour, player experience, affect, retrieval, memory, or commercial-engine performance.

Index Terms—game artificial intelligence, interactive narrative, symbolic validation, constrained generation, runtime enforcement, reproducibility

I. INTRODUCTION

INTERACTIVE generation is more than a language task. A candidate event may be natural yet require an absent item, assert a fact unknown to its speaker, reveal forbidden quest information, introduce an unauthorized effect, or regress a quest stage. Executable text environments make the distinction between a plausible utterance and a valid transition explicit [1], [2]. Grounded dialogue and personalized quest systems further show the value of quest, entity, and graph context [3], [4], but context does not itself authorize canonical mutation.

Structural generation constraints solve an adjacent problem. Grammar-constrained decoding, language-model query languages, and incremental parsers can restrict generated structures [5], [6], [7]. A syntactically admissible event can

nevertheless be false relative to the current world. Retrieval and memory can supply useful evidence [8], [9], [10], but retrieved content is also not a state-transition authority.

This paper presents TRACE-RPG, an implementation that combines these concerns as an auditable transaction protocol. Its current contributions are:

- 1) *versioned trust-boundary contracts* — for canonical state, externally supplied action policy, candidate event, experiment record, and game bridge — enforced by a type-strict parser that rejects ambiguous known-field types and unknown top-level candidate keys; evidenced by deterministic parser-contract fixtures, including proposal/replay rejection parity on a prespecified unknown-key negative;
- 2) *a validate-repair-commit controller* whose six state-relative predicates decide every admitted transition, with bounded repair budget, unchanged deterministic fallback, and defensive pre-commit revalidation; evidenced by exact-count validator/state-isolation and repair-control-flow fixtures over a single authored world;
- 3) *an audit-evidence layer* binding terminal records to unkeyed SHA-256 content checksums, state-semantic replay that recomputes the recorded transition, and episode-continuity validation; evidenced by prespecified fault-injection fixtures in which every named corruption is rejected by its designated check operation;
- 4) *an offline conformance harness* with type-strict known-field adapter boundaries, frozen assignment keys, distinct terminal failure classes, and assignment-complete accounting that retains controller-failure rows without proposal traces; evidenced by manifest-violation fixtures and a frozen release packet with recomputable hashes; and
- 5) *a counterexample-guided repair operator and fixture battery* — $\rho(a, E)$ consumes only the structured error set and the prior candidate, never authoritative state, and is compared against a blind unchanged-retry arm and the state-reading oracle upper bound on an expanded deterministic fixture set partitioned into guided-repairable and guided-irreparable classes; evidenced by exact per-arm commit/fallback counts under frozen manifests.

The empirical center is a deterministic designed-fixture conformance pilot over a single authored world state; a retained authored Godot 4.7.1 policy-mirror trace and render replay add bounded engine-local evidence. Headless runs of that

fixture do record engine-local timing, and we report it as such: each run is one scripted fixture at a fixed seed, and its per-request figure measures local fixture handling rather than model or network inference. Across the four fixture runs, the largest sampled headless frame delta was the first of five samples in every run (98.760–116.667 ms), so the 16.7 ms frame budget was not met; these startup-sensitive traces are not a stable frame benchmark. There are no live model calls, participants, affect or retrieval experiments, live Python–Godot authorization round trips, or representative end-to-end gameplay performance measurements. Those studies are outside this paper.

II. RELATED WORK AND RESEARCH GAP

A. Interactive Narrative and Game Agents

KNUDGE evaluates branching NPC dialogue grounded in quest and entity specifications [3]. Knowledge-graph-assisted quest generation [4], memory–reflection–planning agents [11], and trainable role-playing agents [12] address complementary aspects of grounded and believable behavior. Player-facing studies evaluate preference and perceived dialogue quality [13], [14]. Yin *et al.*’s 2026 online record has an issue assignment that was future-dated at our metadata audit and must be rechecked at submission [15]. A 2025 preprint reports role-sensitive stability–improvisation trade-offs [16]; we use it only as recent motivation. These works motivate separating hard admissibility from optional naturalness and player-experience constructs.

Experience-driven content generation [17] and observer-labelled engagement [18] motivate a future adaptation track. A recent preprint reports limitations in vision-language-model engagement inference [19]. Accordingly, affect remains non-authoritative and is absent from the present implementation evidence.

B. Environments and Benchmarks

TextWorld and ScienceWorld expose explicit interactive state [1], [2]; MineDojo adds open-ended, knowledge-rich embodied tasks [20]. General Video Game AI separates agent, game, and content tracks [21]. AgentBench studies agents across heterogeneous environments [22], AgentBoard adds process-level diagnostics [23], SOTOPIA evaluates social interaction [24], and BALROG evaluates long-horizon language and visual game reasoning [25]. Automated play-style agents provide a complementary testing precedent, not a substitute for human evaluation [26]. These benchmarks inform our future evaluation design but do not prescribe a transactional authorization boundary.

C. Constraints, Neuro-Symbolic Validation, and Retrieval

Constrained decoders improve output structure [5], [6], [7]; TRACE-RPG instead treats structure as an input condition to independent state-relative authorization. Two direct game-specific comparators span complementary boundaries. PANGeA combines memory, validation, a REST service, and Unity integration, but its documented validation asks an LLM

to judge and iteratively refine content against designer rules [27]. IVIE, currently a preprint associated with non-archival ICCC 2026 material, incrementally generates interactive-fiction worlds with symbolic validation [28]. A 2026 system in the same lineage restricts LLMs to selecting authored world-state transformations and reports an exploratory eight-participant study of player expression [29]; it maintains consistency by construction but provides no transactional gate, audit-linked records, or repair-operator evidence. Neither comparison implies superiority: each reinforces that validity is only validity for the objectives and predicates its system encodes.

G-Retriever, HippoRAG, and Mem0 address relevant-subgraph retrieval and long-term memory [8], [9], [10]; Generative Agents and Character-LLM address behavioral memory and persona [11], [12]. TRACE-RPG may later consume such context, but denies it direct commit authority. Retrieval and temporal-memory benefits are not measured here.

D. Auditability and Reproducibility

Process metrics motivate retaining intermediate failures [23]; robust evaluation cautions against overinterpreting sparse stochastic runs [30]. Reproducibility and environmental reporting motivate frozen artifacts and explicit measurement boundaries [31], [32]. Human and LLM evaluation require construct, rater, and bias controls [33], [34], [35], [36], while future non-inferiority and participant analyses require justified margins and random-effects structures [37], [38].

The bounded novelty is the implemented combination, not the invention of its individual parts. Action interposition is long established. Narrative mediation intercepts a user or agent action before it executes and either blocks it or accommodates it by revising the plan [39]; classical planning defines applicability through preconditions and effects [40]; runtime enforcement wraps an untrusted actor in a monitor that admits only permitted transitions [41]; authored-logic interactive drama governs which utterances a character may make, in this journal’s predecessor title [42]; and belief-aware narrative planning reasons over what characters know, which is the concern our knowledge and disclosure predicates encode far more coarsely [43]. TRACE-RPG does not claim to have invented interposition. Its narrower contribution is the evidence layer bound to that gate under an untrusted generative proposer: terminal records linked to content checksums, state-semantic replay of the recorded transition, episode-continuity validation, and assignment-complete failure accounting that retains cases for which no proposal trace can exist. We do not infer that a feature absent from a report is absent from every implementation.

III. PROBLEM DEFINITION AND GUARANTEE BOUNDARY

For exposition, the controller input used by validation at step t is decomposed as

$$c_t = (G_t, q_t), \quad (1)$$

where G_t denotes typed snapshot fields and q_t denotes the externally authored action-policy, disclosure, and quest-stage constraints stored with that snapshot. Terminal record history

$h_{<t}$ is maintained separately from the `WorldState` snapshot; no validation predicate reads it, and the implementation does not reconstruct state from the log. Retrieval, narrative scores, model rationales, memory summaries, and any future affect estimate are non-authoritative proposal context.

An untrusted proposer emits candidate a_t . Strict parsing first enforces the shared candidate-key contract, including rejection of unknown top-level fields, and establishes the declared schema. The validator returns

$$V(c_t, a_t) = (v_{\text{policy}}, v_{\text{pre}}, v_{\text{reach}}, v_{\text{know}}, v_{\text{disc}}, v_{\text{quest}}, E_t), \quad (2)$$

where E_t is a structured error set. Mutation is

$$c_{t+1} = \begin{cases} T(c_t, a_t), & \bigwedge_i v_i = 1, \\ c_t, & \text{otherwise.} \end{cases} \quad (3)$$

The controller asserts the following implementation invariants under its encoded contracts; the pilot tests their mechanisms with deterministic valid, invalid, repair, fallback, replay, and fault-injection fixtures rather than presenting general theorems. **I1 (exhausted-failure immutability)**: if no returned candidate through budget K validates, or proposal/controller execution fails before commit, the returned state equals the supplied prior state. **I2 (bounded recorded attempts)**: provided every proposal and repair callback returns, repair budget $K \geq 0$ admits at most $K + 1$ recorded candidate attempts; a successful commit adds one defensive validation immediately before mutation. The runtime does not enforce callback deadlines, so I2 is not a wall-clock termination guarantee. **I3 (effect authorization)**: an accepted candidate contains no declared fact effect in `CandidateAction.effects` outside `ActionPolicy.allowed_effects`, and any quest-stage target is explicitly authorized by that policy. **I4 (state-semantic replay)**: with the frozen schemas and deterministic transitions, an accepted trace reconstructs its stated terminal state.

These asserted invariants assume correct policy authorship, complete and faithful extraction into declared candidate fields, deterministic software versions, returning callbacks, and single-process file ownership. State-semantic replay revalidates the recorded structured candidate and transition; it does *not* reproduce or verify the provenance, reasoning, prompts, or stochastic generation of intermediate repairs. Natural-language implications omitted from structured fields remain outside the guarantee.

IV. TRACE-RPG ARCHITECTURE AND ALGORITHMS

A. Trust Boundary and Contracts

Fig. 1 places the model or recorded adapter outside the canonical owner. The parser rejects ambiguous numeric and collection types for known fields and rejects unknown top-level candidate fields under the shared key contract. The action policy is supplied separately from the candidate and defines required preconditions and allowed effects. The game bridge is a versioned event interface that may carry observations, proposals, validations, rejections, fallbacks, or commits; only a valid commit authorizes canonical mutation. This decouples interfaces but does not yet demonstrate portability to a




Fig. 1. Trust boundary. Retrieved context and generated candidates remain non-authoritative until known-field type checks and state-relative validation complete.

TABLE I
ENCODED VALIDATOR BOUNDARY

Predicate	Rejects	On failure
Action policy	unknown action/type	unchanged
Precondition	absent/false requirement	unchanged
Reachability	inaccessible object	unchanged
NPC knowledge	undeclared known fact	unchanged
Disclosure	forbidden fact	unchanged
Quest	ineligible/regressive stage	unchanged

commercial engine. Table I summarizes the encoded rejection boundary; it is not a claim of semantic completeness.

B. Validate–Repair–Commit

Fig. 2 gives the controller’s dependency-free pseudocode. “Record” denotes an append to the in-memory trace representation; persistent writing occurs only after terminal consistency checks.

Validation errors are exposed to a repair callback as structured error codes. Repair never directly mutates state, and exhaustion returns an unchanged deterministic fallback. Fig. 3 shows terminal paths. The pilot compares two implemented callbacks. The *counterexample-guided operator* consumes only the prior candidate a and the validator’s structured error set E , never the authoritative state:

$$\rho(a, E) = a \oplus \bigcup_{e \in E} \delta(e), \quad (4)$$

where each $\delta(e)$ is the deterministic per-error-class edit of Table II and $\oplus$ applies the resulting field edits. Each edit touches at most one candidate field per error entity, so the number of changed fields per attempt is bounded by $|E|$; applying ρ twice with the same error set is a fixed point; and an entity simultaneously demanded and rejected on the same field is declared irreparable and left untouched. The *state-reading reference callback* instead discards the error set and reconstructs the policy-governed fields (preconditions, effects, and both quest-stage fields) directly from the authoritative state; it deliberately never edits declaration-integrity fields (required objects, used facts, disclosed facts), because undeclaring a dependency or disclosure would launder the violation past the gate rather than repair the candidate. It is an oracle upper bound over state-readable repairs, and its outcomes must not be read as evidence about any deployable repair method.

```

1:  $a \leftarrow \text{PROPOSE}(c); j \leftarrow 0$ 
2: repeat
3:    $(ok, E) \leftarrow V(c, a); \text{RECORD}(a, E)$ 
4:   if  $ok$  then
5:     assert  $V(c, a).ok$ ; return  $\text{COMMIT}(c, a)$ 
6:   if  $j = K$  then return  $\text{FALLBACK}(c)$ 
7:    $a \leftarrow \text{REPAIR}(c, a, E); j \leftarrow j + 1$ 
8: until terminal

```

Fig. 2. Validate–repair–commit procedure. At most $K + 1$ candidate attempts are recorded; the successful path performs a defensive validation before mutation.

TABLE II
PER-ERROR-CLASS EDITS OF THE GUIDED OPERATOR ρ

Validator error class	Edit $\delta(e)$ on the candidate
Policy precondition omission	insert the named predicate
Policy effect omission	insert the named effect
Policy effect violation	drop the named effect
Unsatisfied precondition	drop the named predicate
Unauthorized stage mutation	clear <code>quest_stage_effect</code>
Unknown action type	no-op (no registered alternative)
Unreachable object; knowledge, disclosure, and stage classes	no-op (candidate-local edit would launder or needs state)

C. Integrity, Semantic Replay, and Accounting

Each result row carries an unkeyed SHA-256 *integrity checksum*; proposal outcomes additionally link to a trace checksum. These hashes detect accidental or tested content changes but are not signatures, message-authentication codes, writer identity, or protection against an actor able to rewrite content and checksums. Semantic replay goes beyond byte integrity by strictly parsing the prior state and every recorded candidate/validation pair, requiring every nonfinal attempt to be invalid, recomputing the terminal transition, and comparing the stated final state. Episode replay also requires state continuity between adjacent JSONL records. Replay does not authenticate or re-execute the repair-generation step that produced candidate $j + 1$.

The terminal record classes are commit, symbolic fallback, adapter failure, and controller failure. A controller-failure row remains in the frozen assignment denominator and leaves state unchanged, but it may legitimately lack a completed proposal trace because no candidate outcome was produced. Accordingly, result completeness does not imply trace existence for every terminal row. The pre-execution assignment key contains run, arm, scenario, seed, model, revision, controller-configuration hash, assignment-input hash, and prior-state hash. Aggregation rejects duplicate, missing, or unexpected keys and separately reports symbolic, adapter, and controller failures.

V. DETERMINISTIC OFFLINE PILOT METHOD

A. Questions and Designed Fixtures

The pilot asks whether: (PQ1) valid candidates commit while named invalid-policy cases remain unchanged; (PQ2) rejection-only, blind unchanged retry, the counterexample-guided operator ρ , and the state-reading reference callback produce the expected outcomes on authored cases in every frozen repairability class; (PQ3) frozen checksum, state-semantic, linkage, and continuity mutations are rejected by




Fig. 3. Proposal, validation, bounded repair, fallback, commit, and replay states. Controller failure is terminal and may have no completed proposal trace.

their designated check operations, while an injected second-file append failure triggers the expected pair-write rollback check; (PQ4) strict adapters and assignment manifests retain exactly one classified terminal row per assigned case or fail closed; and (PQ5) proposal and replay parsers both reject a complete otherwise-valid candidate carrying one unknown top-level key.

Cases are purposively designed conformance fixtures, not samples from a model or game population. The four domains are validator/state isolation, repair-arm comparison, record/trace fault injection, and adapter/accounting contracts. Every repair fixture declares a frozen repairability class before execution: *guided-repairable* (the structured error payload alone carries sufficient information for ρ), *oracle-only* (repair additionally requires reading the authoritative state), or *ir-repairable* (no implemented arm may commit), and each class's expected outcome per arm is part of the designed-fixture contract. Unique cases, rather than repeated deterministic seeds, form descriptive denominators. All expectations are specified before execution.

B. Measures and Reproducibility

Measures are exact counts: expected/observed validator-class agreement; rejected or controller-failure cases with state mutation; commit and attempt outcomes by repair arm and repairability; rejected/prespecified detectable faults by designated operation; replayed/committed traces; rejected/injected manifest violations; and proposal/replay rejection agreement on the prespecified unknown-key negative regression. The last measure is deterministic parser-contract parity, not an empirical safety outcome. Exception class alone does not identify a stable detector layer, so no detector-attribution result is reported. No p -values or population confidence intervals are attached to these designed fixtures. The adapter cases carry synthetic latency and token constants that the input schema pins with JSON-Schema `const`; they exercise accounting-field propagation only and are not measurements of any provider, runtime, or device.

The release packet freezes fixture and assignment manifests, schemas, a result JSON artifact with embedded outcome and trace records, environment and code revision, generated

tables, and checksums. Reporting follows reproducibility and measurement-boundary guidance [31], [32].

VI. PILOT RESULTS

A. Declared-Field Gate Conformance

The pilot loader admits a fixture set only when it isolates every implemented validator code exactly once alongside one valid control, so the 13/13 agreement and the observation of all 12 implemented error codes are *construction invariants of the harness*, not measured success rates. What the run adds is the part the loader does not enforce: each of the 12 isolated negative fixtures showed observed agreement with its authored expected code rather than reaching a different one, and every noncommit path preserved the complete canonical state. This is implementation conformance to an authored structured-field oracle over a single world state, not accuracy against independent semantic labels.

B. Repair-Arm Comparison on Designed Fixtures

The 12 initially invalid designed cases are partitioned into frozen repairability classes: 5 guided-repairable, 1 oracle-only, and 6 irreparable. Rejection-only and blind unchanged-candidate retry each committed 0/12. The counterexample-guided operator ρ committed 5/5 guided-repairable cases and, by design, 0/1 oracle-only and 0/6 irreparable cases; every ρ commit edited a mean of 1.0 candidate fields, and every failed arm execution left the canonical state unchanged. The state-reading reference callback committed 5/5 guided-repairable and 1/1 oracle-only cases, with 0/6 irreparable. These exact counts separate the three mechanisms on the designed fixtures: blind retry never converts a counterexample into a commit, ρ converts exactly the cases whose error payload carries sufficient information, and the oracle additionally converts the case requiring state knowledge. This is operator conformance on frozen fixtures, not repair quality of any live model.

C. Integrity, Boundaries, and Assignment Accounting

All 10 faults prespecified as detectable were rejected by their designated check operations (10/10). Because this pilot does not compare stable typed detector codes, it does not establish detector-layer attribution. Separately, the known provenance-boundary fixture substituted a different invalid precursor before the same recorded valid repair, recomputed the checksum, and passed replay as expected (1/1 undetected). Thus, the mechanism does not protect against a writer who can recompute unkeyed checksums, and replay does not authenticate the repair-generation operation. Two open boundary sentinels—unextracted narrative disclosure and simultaneous omission of a required object from candidate and policy—were intentionally accepted at the encoded layer (2/2), with 0 labelled as safety passes. The former unknown-top-level-field acceptance boundary was reclassified after Stage 8 as a closed negative regression: both proposal and replay parsers specifically rejected a complete 12-field candidate carrying one unknown key (1/1). This establishes parser rejection parity,

TABLE III
REPAIR ARM $\times$ REPAIRABILITY CLASS, EXACT COUNTS (DESIGNED FIXTURES)

Arm	Info source	G-rep. ^f	O-only ^f	Irrep. ^f	Edits ^e
Rejection-only ($K=0$)	none	0/5	0/1	0/6	—
Blind unchanged retry	none	0/5	0/1	0/6	—
Counterexample-guided ρ	error set E	5/5	0/1	0/6	1.0
State-reading oracle	state+policy	5/5	1/1	0/6	1.0

TABLE IV
FROZEN-PILOT RAW ACCOUNTING

Check	Numerator	Denominator
Gate fixture agreement ^c	13	13
Known-boundary acceptances ^a	2	2
Closed unknown-key boundary rejected ^d	1	1
Detectable integrity faults rejected	10	10
Known integrity boundary replay-accepted ^b	1	1
Adapter commits	1	7
Symbolic fallbacks	1	7
Adapter failures	5	7
Manifest guards rejected	3	3

^aBoundary acceptances document semantic-extraction and policy-completeness limits; they are not safety passes. ^bReplay does not authenticate or re-execute repair generation. ^cThe loader enforces this agreement, so it is a construction invariant rather than a measurement. ^dThis is post-Stage-8 parser rejection parity, not semantic-safety evidence. ^eMean changed candidate fields over committed repairs (minimality proxy); failed arm executions left the state unchanged. ^fCommits/cases. G-rep. = guided-repairable, O-only = oracle-only, Irrep. = irreparable (frozen repairability classes).

not semantic safety. Among seven adapter/accounting assignments, outcomes were 1 commit, 1 symbolic fallback, and 5 adapter failures. Synthetic telemetry fields were populated and propagated for 2/7 assignments; these are schema-pinned constants verifying accounting-field propagation, not measurements. All 3/3 injected duplicate or missing-assignment guards failed closed. These are raw counts from synthetic offline fixtures.

Tables III and IV report exact designed-fixture counts: the repair table separates the four arms per frozen repairability class, and neither table supports a population or live-model estimate.

Fig. 5 makes the reporting boundary explicit: only rows admitted by the frozen fixture and designated check operation enter the generated result tables.

Godot 4.7.1 replayed the retained *Sealed Lighthouse* canonical trace in a non-headless 1280 $\times$ 720 SubViewport. Its manifest binds each PNG to event/delivery indices, validation, and state hashes: Fig. 4(b) preserves the before hash, whereas (c) changes it after a valid commit. All source panels are generated-asset-free primary-track renders; the illustrative crops below enlarge only their authored status/event callouts.

The selected v5 packet establishes one scripted fixture’s render/state correspondence only; it adds no live model or Python transport, participant or player input, adaptation intervention, or narrative-quality result.

OBSERVATION		FALLBACK · NO MUTATION		COMMIT · REPLAY-BOUND	
EVENT evt-000-observe		EVENT evt-002-fallback-secret		EVENT evt-005-commit-hint	
(a) Arrival observation		(b) Unchanged-state rejection		(c) Authorized commit	
Panel	Capture	Event	Validation	State relation	
(a)	sl-rc-001-arrival	evt-000-observe	not applicable	snapshot = before	
(b)	sl-rc-002-rejected-secret	evt-002-fallback-secret	invalid	world = before	
(c)	sl-rc-003-authorized-hint	evt-005-commit-hint	valid	world $\neq$ before	

Full-resolution artifact pointer: godot-4.7.1-20260813t115916z-sealed-lighthouse-render-v5/rendered-canonical-v1/; filenames are the capture IDs above plus .png.

Fig. 4. Illustrative status/event crops from the authored Godot render replay, with manifest metadata. The trace-bound correspondence is not live-transport, performance, usability, immersion, or efficacy evidence.




Fig. 5. Evidence boundary from frozen fixture through a designated check operation and generated table. Counts are admitted only through generated artifacts.

VII. DISCUSSION

A. Established Scope

The designed pilot can establish mechanism behavior only for its frozen cases: whether encoded checks isolate state, whether bounded repair and fallback follow their declared paths, whether named corruptions are rejected by their designated operations, and whether assigned rows remain classified for aggregation. It does not attribute a rejection to a stable typed detector code. This is useful for authoring-time validation, continuous-integration regression, recorded-adapter tests, and pre-engine replay inspection.

The guided-versus-oracle boundary is the honest reading of the repair comparison. The state-reading callback remains the upper bound: it may consult the authoritative state, so it converts the oracle-only case that ρ , by construction, cannot. The counterexample-guided operator matches the oracle exactly on the guided-repairable class while consuming only the validator’s structured counterexamples—the same information any deployed proposer would receive at the trust boundary. That separation, not a success rate, is the finding: it locates which repairs are reachable from error payloads alone and which structurally require state access, on these designed fixtures. Whether a live model prompted with the same error payloads approaches the ρ row is an open empirical question outside this paper. Feedback-consuming repair loops are precedented: Self-Refine feeds back language generated by the same untrusted model [44], and Self-Debugging consumes execution results from a trusted full-program runtime [45]. ρ differs in both

channel and evidence class: its feedback is a typed error set computed by the deterministic validator before any execution, and it is measured as an operator on designed fixtures, not as a live-model capability.

B. Relationship to Prior Work

TRACE-RPG complements constrained decoders by applying state-relative authorization after parsing [5], [6], [7]. It differs from retrieval and memory components by withholding commit authority [8], [9], [10]. It complements executable environments and benchmarks rather than replacing them [1], [2], [22], [25]. Process-linked traces instantiate a measurement concern emphasized by AgentBoard [23]; they are not evidence that TRACE-RPG outperforms that or any other system.

The studies that would test efficacy—a preregistered multi-model comparison with an independently labelled semantic oracle and a genuine blind-retry comparator, retrieval and memory ablations at matched encoded validity, and a blinded matched-validity human study with justified power and mixed-effects analysis—are outside this paper and are not reported here in any partial form [22], [33], [30], [38].

VIII. THREATS, ETHICS, AND ARTIFACT SCOPE

The implementation evidence is bounded by nine explicit limitations. **L1 (encoded-semantics boundary):** validation rejects only what declared fields and predicates represent; natural-language implications omitted from structured fields pass untested, and an encoded check cannot discover semantics its authors did not encode. **L2 (single world):** all counts are conditioned on one authored world state and policy; cross-world, cross-policy, and scale generalization are untested by construction. **L3 (repair evidence boundary):** the state-reading reference callback remains an oracle upper bound; the counterexample-guided operator $\rho(a, E)$ is evidenced only by deterministic fixture counts over guided-repairable and guided-irreparable classes, and neither result measures live-model repair quality. **L4 (no live transport):** proposers and adapters are recorded or synthetic; no live model calls or live engine authorization round trips exist, and adapter telemetry is schema-pinned constants, not measurement. **L5 (no human data):** no participants or affect measurements exist; player-facing legibility of the commit ritual is a design hypothesis,

not an evaluated outcome. **L6 (integrity, not authentication):** unkeyed SHA-256 checksums detect accidental or tested corruption; they provide no writer identity or adversarial tamper resistance. **L7 (single-process writer):** record consistency assumes single-process file ownership; concurrent-writer safety is neither implemented nor tested. **L8 (engine portability undemonstrated):** the versioned bridge decouples interfaces, but engine evidence is one scripted Godot fixture with startup-dominated frame samples that miss the 16.7ms budget; no commercial-engine or real-time result exists. **L9 (designed fixtures, no population inference):** fixtures are purposively authored conformance probes; exact counts support implementation conformance only and carry no population or safety-rate inference. Internal validity additionally depends on independently authored policy expectations and absence of fixture leakage, and reproducibility on fixed software and artifact versions.

The pilot processes no participants or personal telemetry. Future human and affect work requires applicable review, informed consent, compensation disclosure, minimization, retention/deletion rules, and opt-out. LLM judges, if used, remain auxiliary because both automated and crowd judgements require calibration and bias controls [33], [34], [35], [36].

IX. CONCLUSION

TRACE-RPG isolates authored candidate-event payloads from canonical mutation through an encoded symbolic commit gate and makes terminal outcomes auditable through linked records and state-semantic replay. The deterministic offline pilot is evidence about the implemented mechanism in frozen fixture domains over a single authored world state, not cross-model superiority, player-experience benefit, retrieval or memory effectiveness, affect efficacy, or commercial-engine performance. These remain explicit confirmatory extensions.

DISCLOSURE OF AI USAGE

Large language models assisted prose, code, and test drafting; command orchestration; evidence auditing and result interpretation; the reviewer-panel simulation that drove this revision; and citation-identity checks against bibliographic metadata APIs. They did not replace deterministic execution as the source of reported values. Pilot counts are emitted by the project runner from authored fixtures and enter the manuscript through machine-generated fragments; Godot timing and render statements are transcribed from the retained engine-evidence packet. Implementation and result claims in this revision were cross-checked against source and artifacts, and every citation identity was checked against at least one independent index. The authors take full responsibility for the content.

DATA AND CODE AVAILABILITY

The repository contains the implementation, pilot bundle, fixture and assignment manifests, generated result fragments, and build sources. The current pilot SHA-256 manifest covers 38 artifacts and 22 declared inputs and binds 121 provenance rows (85 executed fixture rows and

36 aggregate rows); all artifact and input hashes recompute from the frozen packet. The manifest records the portable invocation `uv run python` and contains no absolute user or clone path. The clean-Git input binding (`dirty=false` at a tagged input commit, previously closed at tag `trace-rpg-stage10-inputs-20260814-v1`) is re-established at tag `trace-rpg-guided-repair-inputs-20260821-v1` in this release; it is a release gate, not a property of intermediate working trees. None of this means that an anonymized reviewer archive or DOI-bearing deposit already exists, and neither is claimed here.

REFERENCES

- [1] M.-A. Côté, Á. Kádár, X. Yuan *et al.*, “TextWorld: A learning environment for text-based games,” in *Computer Games*, ser. Communications in Computer and Information Science. Springer International Publishing, 2019, vol. 1017, pp. 41–75.
- [2] R. Wang, P. Jansen, M.-A. Côté, and P. Ammanabrolu, “ScienceWorld: Is your agent smarter than a 5th grader?” in *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2022, pp. 11 279–11 298.
- [3] N. Weir, R. Thomas, R. d’Amore, K. Hill, B. Van Durme, and H. Jhamtani, “Ontologically faithful generation of non-player character dialogues,” in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2024, pp. 9212–9242.
- [4] T. Ashby, B. K. Webb, G. Knapp, J. Searle, and N. Fulda, “Personalized quest and dialogue generation in role-playing games: A knowledge graph- and language model-based approach,” in *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*. ACM, 2023, pp. 1–20.
- [5] S. Geng, M. Josifoski, M. Peyrard, and R. West, “Grammar-constrained decoding for structured NLP tasks without finetuning,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2023, pp. 10932–10952.
- [6] L. Beurer-Kellner, M. Fischer, and M. Vechev, “Prompting is programming: A query language for large language models,” *Proceedings of the ACM on Programming Languages*, vol. 7, no. PLDI, pp. 1946–1969, 2023.
- [7] T. Scholak, N. Schucher, and D. Bahdanau, “PICARD: Parsing incrementally for constrained auto-regressive decoding from language models,” in *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2021, pp. 9895–9901.
- [8] X. He, Y. Tian, Y. Sun *et al.*, “G-Retriever: Retrieval-augmented generation for textual graph understanding and question answering,” in *Advances in Neural Information Processing Systems*, vol. 37. Neural Information Processing Systems Foundation, 2024, pp. 132 876–132 907.
- [9] B. Jiménez Gutiérrez, Y. Shu, Y. Gu, M. Yasunaga, and Y. Su, “HippoRAG: Neurobiologically inspired long-term memory for large language models,” in *Advances in Neural Information Processing Systems*, vol. 37. Neural Information Processing Systems Foundation, 2024, pp. 59 532–59 569.
- [10] P. Chhikara, D. Khant, S. Aryan, T. Singh, and D. Yadav, “Mem0: Building production-ready AI agents with scalable long-term memory,” in *ECAI 2025*, ser. Frontiers in Artificial Intelligence and Applications. IOS Press, 2025, peer-reviewed ECAI 2025 record; page range was not exposed by the accessible Crossref or publisher metadata at audit time.
- [11] J. S. Park, J. C. O’Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein, “Generative agents: Interactive simulacra of human behavior,” in *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*. ACM, 2023, pp. 1–22.
- [12] Y. Shao, L. Li, J. Dai, and X. Qiu, “Character-LLM: A trainable agent for role-playing,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2023, pp. 13 153–13 187.
- [13] N. Akoury, Q. Yang, and M. Iyyer, “A framework for exploring player perceptions of LLM-generated dialogue in commercial video games,” in *Findings of the Association for Computational Linguistics: EMNLP 2023*. Association for Computational Linguistics, 2023, pp. 2295–2311.
- [14] M. Hochreiter, S. Kriglstein, and G. Wallner, “Beyond pre-defined scripts: Player perceptions on generative non-player character dialogues,” in *Proceedings of the 31st International Conference on Intelligent User Interfaces*. ACM, 2026, pp. 2004–2018.
- [15] M. Yin, H. Zu, W. Gu *et al.*, “How contextualized generative AI shapes player experience in games,” *Entertainment Computing*, vol. 58, p. 101194, 2026, online record with a future issue assignment at the 2026-08-12 audit; issue metadata must be rechecked at submission.
- [16] V. Figueiredo and D. Elumeze, “Symbolically scaffolded play: Designing role-sensitive prompts for generative NPC dialogue,” 2025, preprint; no archival venue or publisher DOI verified as of 2026-08-12.
- [17] G. N. Yannakakis and J. Togelius, “Experience-driven procedural content generation,” *IEEE Transactions on Affective Computing*, vol. 2, no. 3, pp. 147–161, 2011.
- [18] D. Melhart, D. Gravina, and G. N. Yannakakis, “Moment-to-moment engagement prediction through the eyes of the observer: PUBG streaming on Twitch,” in *International Conference on the Foundations of Digital Games*. ACM, 2020, pp. 1–10.

- [19] Z. Wang, Q. Guo, R. Singh, and X. Hu, “Do vision language models understand human engagement in games?” 2026, preprint; no archival venue or publisher DOI verified as of 2026-08-12.
- [20] L. Fan, G. Wang, Y. Jiang *et al.*, “MineDojo: Building open-ended embodied agents with internet-scale knowledge,” in *Advances in Neural Information Processing Systems*, vol. 35. Neural Information Processing Systems Foundation, 2022, pp. 18 343–18 362.
- [21] D. Pérez-Liébana, J. Liu, A. Khalifa, R. D. Gaina, J. Togelius, and S. M. Lucas, “General video game AI: A multitrack framework for evaluating agents, games, and content generation algorithms,” *IEEE Transactions on Games*, vol. 11, no. 3, pp. 195–214, 2019.
- [22] X. Liu, H. Yu, H. Zhang *et al.*, “AgentBench: Evaluating LLMs as agents,” in *International Conference on Learning Representations*, 2024, pp. 52 989–53 046, peer-reviewed official proceedings record; no publisher DOI assigned.
- [23] C. Ma, J. Zhang, Z. Zhu *et al.*, “AgentBoard: An analytical evaluation board of multi-turn LLM agents,” in *Advances in Neural Information Processing Systems*, vol. 37. Neural Information Processing Systems Foundation, 2024, pp. 74 325–74 362.
- [24] X. Zhou, H. Zhu, L. Mathur *et al.*, “SOTOPIA: Interactive evaluation for social intelligence in language agents,” in *International Conference on Learning Representations*, 2024, pp. 40 975–41 019, peer-reviewed official proceedings record; no publisher DOI assigned.
- [25] D. Paglieri, B. Cupiał, S. Coward *et al.*, “BALROG: Benchmarking agentic LLM and VLM reasoning on games,” in *International Conference on Learning Representations*, 2025, pp. 96 666–96 702, peer-reviewed official proceedings record; no publisher DOI assigned.
- [26] P. Le Pelletier de Woillemont, R. Labory, and V. Corruble, “Automated play-testing through RL-based human-like play-styles generation,” *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, vol. 18, no. 1, pp. 146–154, 2022.
- [27] S. Buongiorno, L. Klinkert, Z. Zhuang, T. Chawla, and C. Clark, “PANGeA: Procedural artificial narrative using generative AI for turn-based, role-playing video games,” *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, vol. 20, no. 1, pp. 156–166, 2024.
- [28] M. Vaucher, S. Silveira, S. Góngora, and L. Chiruzzo, “IVIE: A neuro-symbolic approach to incremental and validated generation of interactive fiction worlds,” 2026, preprint; accepted/forthcoming at ICCV 2026, whose available pre-proceedings were explicitly non-archival and unpublished at retrieval time.
- [29] S. Góngora, L. Chiruzzo, G. Méndez, and P. Gervás, “World-state transformations for neuro-symbolic interactive storytelling,” 2026, preprint; no archival venue or publisher DOI verified as of 2026-08-21.
- [30] R. Agarwal, M. Schwarzer, P. S. Castro, A. Courville, and M. G. Bellemare, “Deep reinforcement learning at the edge of the statistical precipice,” in *Advances in Neural Information Processing Systems*, vol. 34. Curran Associates, Inc., 2021, pp. 29 304–29 320, peer-reviewed official proceedings record; no publisher DOI assigned.
- [31] J. Pineau, P. Vincent-Lamarre, K. Sinha *et al.*, “Improving reproducibility in machine learning research: A report from the NeurIPS 2019 reproducibility program,” *Journal of Machine Learning Research*, vol. 22, no. 164, pp. 1–20, 2021, peer-reviewed official journal article; no publisher DOI assigned.
- [32] P. Henderson, J. Hu, J. Romoff, E. Brunskill, D. Jurafsky, and J. Pineau, “Towards the systematic reporting of the energy and carbon footprints of machine learning,” *Journal of Machine Learning Research*, vol. 21, no. 248, pp. 1–43, 2020, peer-reviewed official journal article; no publisher DOI assigned.
- [33] C. van der Lee, A. Gatt, E. van Miltenburg, S. Wubben, and E. Krahmer, “Best practices for the human evaluation of automatically generated text,” in *Proceedings of the 12th International Conference on Natural Language Generation*. Association for Computational Linguistics, 2019, pp. 355–368.
- [34] M. Karpinska, N. Akoury, and M. Iyyer, “The perils of using Mechanical Turk to evaluate open-ended text generation,” in *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2021, pp. 1265–1285.
- [35] L. Zheng, W.-L. Chiang, Y. Sheng *et al.*, “Judging LLM-as-a-judge with MT-Bench and Chatbot Arena,” in *Advances in Neural Information Processing Systems*, vol. 36. Neural Information Processing Systems Foundation, 2023, pp. 46 595–46 623.
- [36] G. H. Chen, S. Chen, Z. Liu, F. Jiang, and B. Wang, “Humans or LLMs as the judge? a study on judgement bias,” in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2024, pp. 8301–8327.
- [37] D. Lakens, A. M. Scheel, and P. M. Isager, “Equivalence testing for psychological research: A tutorial,” *Advances in Methods and Practices in Psychological Science*, vol. 1, no. 2, pp. 259–269, 2018.
- [38] D. J. Barr, R. Levy, C. Scheepers, and H. J. Tily, “Random effects structure for confirmatory hypothesis testing: Keep it maximal,” *Journal of Memory and Language*, vol. 68, no. 3, pp. 255–278, 2013.
- [39] M. O. Riedl, C. J. Saretto, and R. M. Young, “Managing interaction between users and agents in a multi-agent storytelling environment,” in *Proceedings of the Second International Joint Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, 2003, pp. 741–748.
- [40] R. E. Fikes and N. J. Nilsson, “STRIPS: A new approach to the application of theorem proving to problem solving,” *Artificial Intelligence*, vol. 2, no. 3–4, pp. 189–208, 1971.
- [41] F. B. Schneider, “Enforceable security policies,” *ACM Transactions on Information and System Security*, vol. 3, no. 1, pp. 30–50, 2000.
- [42] R. Evans and E. Short, “Versu—a simulationist storytelling system,” *IEEE Transactions on Computational Intelligence and AI in Games*, vol. 6, no. 2, pp. 113–130, 2014, predecessor title of IEEE Transactions on Games.
- [43] S. G. Ware and C. Siler, “Sabre: A narrative planner supporting intention and deep theory of mind,” in *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment (AIIDE)*, vol. 17, no. 1, 2021, pp. 99–106.
- [44] A. Madaan, N. Tandon, P. Gupta *et al.*, “Self-Refine: Iterative refinement with self-feedback,” 2023, preprint; commonly indexed as NeurIPS 2023 but proceedings identity not re-verified as of 2026-08-21.
- [45] X. Chen, M. Lin, N. Schärli, and D. Zhou, “Teaching large language models to self-debug,” 2023, preprint; no archival venue or publisher DOI verified as of 2026-08-21.